

Problem sheet #01
Erza Dauti

Problem 1.1: *library and system calls*

(1+1 = 2 points)

Answer the following questions by using `strace` and `ltrace` on a Linux system. Provide enough context information to make it clear how the results were obtained.

- a) How many system calls and how many library calls does executing `/bin/date` produce?

- **System calls (using `strace`) :**

To capture the number of system calls, we ran the following command:

```
~$ strace /bin/date
```

This output showed that **118 system calls** were made during the execution of `/bin/date`.

- **Library calls (using `ltrace`) :**

To trace the number of library calls, we used:

```
~$ ltrace /bin/date
```

This showed that **47 library calls** were made during the execution.

- b) What are the most frequent (top three) library and system calls and what do these calls do?

Top Three System Calls (via `strace`):

1. `openat()`: 31 calls
2. `fstat()`: 21 calls
3. `close()`: 21 calls

Top Three Library Calls (via `ltrace`):

1. `fwrite()`: 8 calls
2. `fputc()`: 7 calls
3. `fread()`: 4 calls

Problem 1.2: system call errors

(1+1 = 2 points)

System call errors are usually indicated by returning a special value (usually -1 for system calls that return an int) and by indicating the details in the global variable `int errno`, declared in `errno.h`.

- a) For each of the following system calls, describe a condition that causes it to fail (i.e., a condition that causes -1 to be returned and that sets `errno` to a distinct value).

- `int open(const char *path, int oflag, ...)`

- **[ENXIO]**

The named file is a character special or block special file, and the device associated with this special file does not exist.

- `int close(int fildes)`

- **[EINTR]**

The `close()` function was interrupted by a signal.

- b) What is the value of `errno` after a system call is completed without an error?

- When a system call completes without an error, the value of `errno` remains unchanged.
By default, the value of `errno` is 0, and it will stay at 0 if no error occurs during a system call.

`errno` is only updated when a system call fails, in which case it is assigned a non-zero value that represents a specific error code

Therefore, if a system call completes without errors, **`errno` remains at 0** or retains its previous value if there were no prior errors.

Problem 1.3: *execute a command in a modified environment or print the environment* (6 points)

On Unix systems, processes have access to environment variables that can influence the behavior of programs. The global variable `environ`, declared as

```
extern char **environ;
```

points to an array of pointers to strings. The last pointer has the value `NULL`. By convention, the strings have the form “name=value” and the names are often written using uppercase characters. Examples of environment variables are `USER` (the name of the current user), `HOME` (the current user’s home directory), or `PATH` (the colon-separated list of directories where the system searches for executables).

Write a program `env` that implements some of the functionality of the standard `env` program. The syntax of the command line arguments is the following:

```
env [OPTION]... [NAME=VALUE]... [COMMAND [ARG]...]
```

- a) If called without any arguments, `env` prints the current environment to the standard output.
- b) If called with a sequence of “name=value” pairs and no further arguments, the program adds the “name=value” pairs to the environment and then prints the environment to the standard output.
- c) If called with a command and optional arguments, `env` executes the command with the given arguments.
- d) If called with a sequence of “name=value” pairs followed by a command and optional arguments, the program adds the “name=value” pairs to the environment and executes the command with the given arguments in the modified environment.
- e) If called with the option `-v`, the program writes a trace of what it is doing to the standard error.
- f) If called with the option `-u name`, the program removes the variable `name` from the environment.

Here are some example invocations:

<code>\$ env</code>	<code># print the current environment</code>
<code>\$ env foo=bar</code>	<code># add foo=bar and print the environment</code>
<code>\$ env -u foo</code>	<code># remove foo and print the environment</code>
<code>\$ env date</code>	<code># execute the program date</code>
<code>\$ env TZ=GMT date</code>	<code># add TZ=GMT and execute the program date</code>
<code>\$ env -u TZ date</code>	<code># remove TZ and execute the program date</code>
<code>\$ env -u x a=b b=c date</code>	<code># remove x, add a and b, execute date</code>

Hand in the source code of your `env` program. Make sure that your program handles *all* error situations appropriately. Use the `getopt()` function of the C library for parsing command line options. Furthermore, use one of the `exec` system calls like `execvp()` to execute a command. (Using `system()` can be made to work but it is somewhat difficult to get right since concatenating strings using space characters may lead to surprises if the strings themselves contain space characters; to do this correctly, you have to quote the strings such that the shell called by the `system()` library function tokenizes the string properly again. Naive concatenation usually leads to a security weakness, it is often better to avoid the `system()` library function. See also the Caveats section in the Linux manual page describing the `system()` library function.)

```

#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#include <unistd.h>
#include <getopt.h>

extern char **environ;

// Print environment variables
void show_environment() {
    char **env_ptr = environ;
    while (*env_ptr) {
        printf("%s\n", *env_ptr);
        env_ptr++;
    }
}

int main(int arg_total, char *arg_list[]) {
    int flag;
    int remove_flag = 0, verbose_flag = 0;

    opterr = 0;
    while ((flag = getopt(arg_total, arg_list, "vu:")) != -1) {
        switch (flag) {
            case 'v':
                verbose_flag = 1;
                break;
            case 'u':
                if (optarg[0] == '-') {
                    fprintf(stderr, "Invalid -u option argument: '%s'\n", optarg);
                    return EXIT_FAILURE;
                }
                remove_flag = 1;
                if (unsetenv(optarg)) {
                    perror("unsetenv");
                    return EXIT_FAILURE;
                }
                break;
            default:
                fprintf(stderr, "Usage: %s [-v] [-u name] [name=value]... [command [arg...]]!\n", arg_list[0]);
                exit(EXIT_FAILURE);
                break;
        }
    }

    //Call without any arguments
    if (arg_total == 1) {
        show_environment();
    }

    if ((verbose_flag == 1) && (remove_flag == 1)) {
        for (int i = 1; i < optind; i++) {
            if (!strcmp(arg_list[i], "-u")) {
                fprintf(stderr, "Removed %s!\n", arg_list[i+1]);
            }
        }
    }

    // Traverse the arguments that aren't options
    for (int i = optind; i < arg_total; i++) {
        if (strchr(arg_list[i], '=')) {
            char *pair_copy = arg_list[i];
            char *key = strtok(pair_copy, "=");
            char *val = strtok(NULL, "=");

            if ((key == NULL) || (val == NULL)) {
                fprintf(stderr, "Invalid name=value pair!\n");
                return EXIT_FAILURE;
            }

            if (setenv(key, val, 1)) {
                perror("setenv");
                return EXIT_FAILURE;
            }

            if (verbose_flag == 1) {
                fprintf(stderr, "Added %s=%s pair!\n", key, val);
            }
        }
        else {
            // Print the name of the program to be executed
            if (verbose_flag == 1) {
                fprintf(stderr, "Executing %s!\n", arg_list[i]);
            }
            // Run
            execvp(arg_list[i], &arg_list[i]);
            perror("execvp");
            return EXIT_FAILURE;
        }
    }

    if ((remove_flag == 0) && (verbose_flag == 0)) {
        show_environment();
    }
    return EXIT_SUCCESS;
}

```

